


Interactive Computing and Running Jupyter Notebooks

Presented at the Using Python and Jupyter Notebooks on TSCC Workshop

September 2, 2021

Mary Thomas
(SDSC)



Outline

- Defining Interactive High-Performance Computing
- Running Interactive Jobs
- Application Examples

Resources

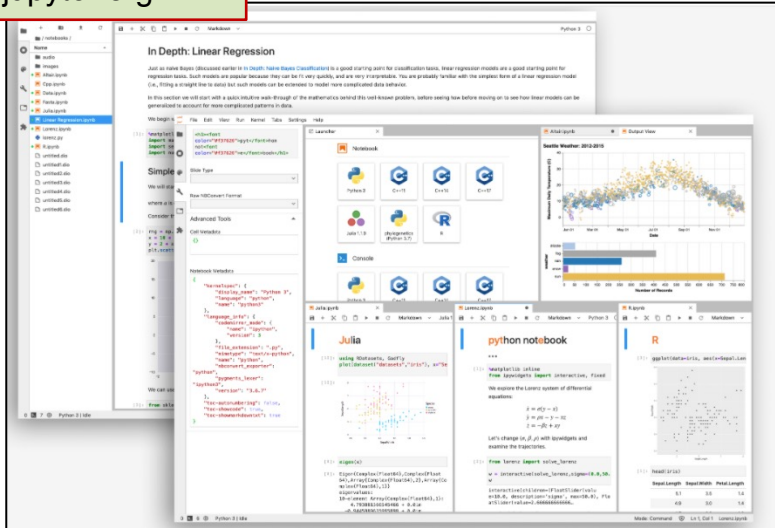
- TSCC User Guide
 - https://www.sdsc.edu/support/user_guides/tsc.html
- GitHub Repo for this workshop:
 - <https://github.com/sdsc-hpc-training-org/Using-Python-and-Jupyter-Notebooks-on-TSCC/>
- GitHub Repo for Example Notebooks:
 - <https://github.com/sdsc-hpc-training-org/notebook-examples>
- SDSC Training Resources
 - https://www.sdsc.edu/education_and_training/training

What is Interactive HPC-Computing

- In **computer** science, **interactive computing** refers to software which accepts input from the user as it runs.
 - commonly used programs, such as word processors or spreadsheet applications.
- **Interactive HPC computing** involves *real-time* user inputs to perform tasks on a set of compute node(s) including:
 - Code development, real-time data exploration, and visualizations.
 - Used when applications have large data sets or are too large to download to local device, software is difficult install, etc.
 - User inputs come via command line interface or application GUI (Jupyter Notebooks, Matlab, R-studio).
 - Actions performed on remote compute nodes as a result of user input or program out.

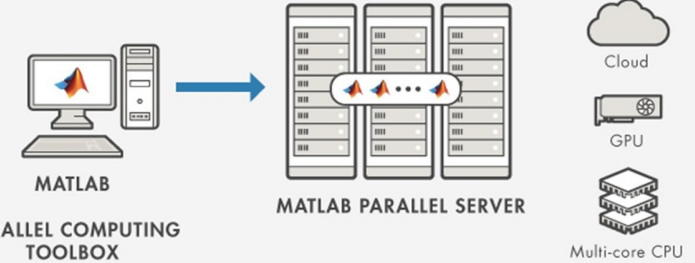
Interactive HPC Scenarios

<https://jupyter.org/>



Parallel Matlab (AWS)

```
>> parpool(parcluster('HPC1'),100);
>> parfor i=1:3000
>>     c(i,:) = eig(rand(1000));
>> end
```




<https://azuremarketplace.microsoft.com/en-us/marketplace/apps/mathworks-inc.matlab-parallel-server-listing?tab=Overview>

Interactive Distributed Computing with Jupyter (NERSC)

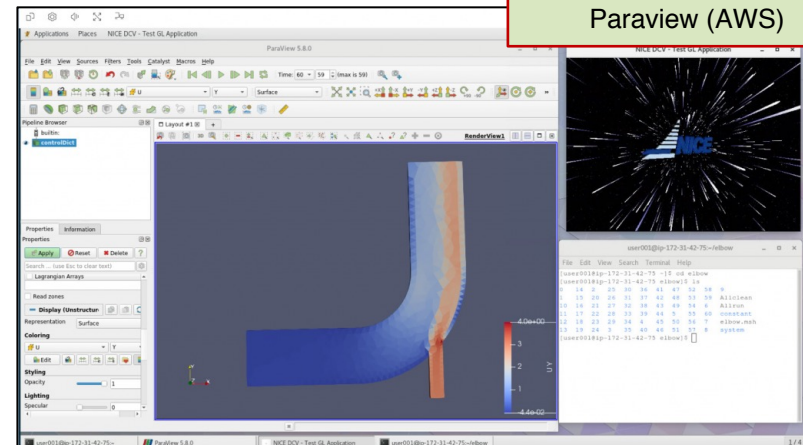
Jupyter architecture

- Allocate nodes on Cori interactive queue and start ipyparallel or Dask cluster
 - Developed %ipcluster magic to setup within notebook
- Compute nodes traditionally do not have external address
 - Required network configuration / policy decisions
- Distributed training communication is via MPI Horovod or Cray ML Plugin



<https://drive.google.com/file/d/1-OFjrkl1q3L1d3uakr2xkoZrPn2c2VZpZ/view>

Paraview (AWS)



<https://aws.amazon.com/blogs/compute/how-to-run-3d-interactive-applications-with-nice-dcv-in-aws-batch/>

TSCC Interactive Jobs

- Must be run on compute nodes (CPU, GPU)
 - Depends on which partition you choose
- Submit a request to the queuing system
- Choose which queue you want to use
- Use the *qsub* command:

```
qsub -I -l nodes=2:ppn=10 -l walltime=0:50:00
```

Partition Name	Max Walltime	Max Processors/User	Max Running + Queued Jobs	Accessible by	Comments
hotel	168 hrs	128	1500	all	Regular compute nodes for all users
gpu-hotel	168 hrs	--	--	all	GPU nodes for all users
pdafm	168 hrs	96	50	all	pdafm (large memory) nodes for all users
home	unlimited	unlimited	1500	condo	Home node(s) for condo participants
condo	8 hrs	512	1500	condo	Compute nodes for condo participants
gpu-condo	8 hrs	84	--	condo	gpu nodes for condo participants
glean	1 hr	1024	500	condo	pre-emptible nodes for condo participant, free of charge

HPC Basics - Running jobs on TSCC

- All jobs submitted via TORQUE resource manager (pbs) with Maui scheduler.
- Batch jobs : submit batch scripts from login nodes.
 - `qsub test.sh`
 - Submit script to include PBS directives such as
 - number of nodes, cores, estimated time
 - right queue
 - job name, email address
 - account
- Interactive jobs:
 - `qsub -I -l nodes=1:ppn=1 -t 00:30:00 -q condo -A account`

Some useful commands...

- `#PBS -k oe =>` writes directly to destination dir
- **qalter** – change parameters of a queued job (e.g. email, account)
- **qdel** <jobnumber> – delete job

Monitoring Job Submission

- qstat - current status of jobs and queues
 - qstat -u username
 - qstat -f jobnumber
- yqd – prints out why the job hasn't started

```
$ yqd 13977625  
13977625 (xxxxxx home-YXX 1x24 ['haswell'] 0:12:23): 0 nodes free
```

- lsjobs - useful to see the node availability

```
$ lsjobs --property=condo-node  
tscc-4-69: FREEx28  
tscc-4-70: FREEx28
```

Job Monitoring contd...

- checkjob

```
$ checkjob 13976617
```

```
.....
```

```
Reservation '13976617' (9:19:34:24 -> 12:19:34:24 Duration:  
3:00:00:00)
```

```
PE: 3.00 StartPriority: 203174
```

```
job cannot run in partition DEFAULT (idle procs do not meet  
requirements
```

```
: 0 of 3 procs found)
```

```
idle procs: 3995 feasible procs: 0
```

```
.....
```

- showq -l (position in queue)
- showstart <jobid> (estimated start time)

Account Monitoring

- Users can belong to multiple groups

groups <username>

- *groups <username>* belonging to multiple groups

- Accounts are used to keep track of allocation usage

gbalance -u <username>

- Specified by #PBS -A
- Check usage with

gstatement -u <username>

TSCC Interactive Jobs: Application Examples

- MPI "Hello World"
- Matlab
- R console
- Jupyter Notebooks → Covered in next two talks

TSCC Interactive Job: Hello-MPI

```
qsub -I -l nodes=2:ppn=10 -l walltime=0:50:00
```

Interactive nodes can be used for running parallel jobs (MPI, OpenMP, CUDA code if GPU), compiling large codes, etc.

```
[mthomas@tscc-gpu-9-4 hello-mpi]$  
[mthomas@tscc-gpu-9-4 hello-mpi]$ qsub -I -l nodes=1:ppn=16 -l walltime=0:30:00 -q glean  
qsub: waiting for job 26842841.tscc-mgr7.local to start  
qsub: job 26842841.tscc-mgr7.local ready
```

[mthomas@tscc-gpu-9-4 ~]\$ pwd
[mthomas@tscc-gpu-9-4 hello-mpi]\$ mpirun -np 8 ./hello-mpi

```
node      0 : Hello and Welcome to TSCC!  
node      6 : Hello and Welcome to TSCC!  
node      1 : Hello and Welcome to TSCC!  
node      3 : Hello and Welcome to TSCC!  
node      4 : Hello and Welcome to TSCC!  
node      5 : Hello and Welcome to TSCC!  
node      7 : Hello and Welcome to TSCC!  
node      2 : Hello and Welcome to TSCC!
```

Request interactive node
using glean queue node

TSCC Interactive Jobs: Applications

Running Matlab without GUI

```
qsub -l -l nodes=1:ppn=16 walltime=0:30:00 -q glean
```

```
[mthomas@tscc-login1 hello-mpi]$ qsub -I -l nodes=1:ppn=16 -l  
walltime=0:30:00 -q glean  
qsub: waiting for job 26842936.tscc-mgr7.local to start  
qsub: job 26842936.tscc-mgr7.local ready
```

```
[mthomas@tscc-13-12 ~]$ module load matlab  
[mthomas@tscc-13-12 ~]$ matlab -nodisplay  
< M A T L A B (R) >  
Copyright 1984-2016 The MathWorks, Inc.  
R2016b (9.1.0.441655) 64-bit (glnxa64)  
September 7, 2016
```

To get started, type one of these: helpwin, helpdesk, or demo.
For product information, visit www.mathworks.com.

```
>> A = [1 3 0; 2 4 -1; 4 9 -1]  
A =  
     1     3     0  
     2     4    -1  
     4     9    -1  
>> B=A'  
B =  
     1     2     4  
     3     4     9  
     0    -1    -1  
>> A*B  
ans =  
    10    14    31  
    14    21    45  
    31    45    98  
>> exit  
[mthomas@tscc-13-12 ~]$
```

TSCC Interactive Jobs: Running Matlab with GUI

To use a Graphical User Interface (GUI) as part of your interactive job, **you will need to set up Xforwarding**. Example below is for using XQuartz on a MAC. For examples for MacOS and Windows, see: <http://systems.eecs.tufts.edu/x11-forwarding/>

Step 1: Set X11 forwarding on the computer that you are connecting from: Update or install Xquartz and **restart** your system

```
[mthomas@home]$ brew install --cask xquartz
Updating Homebrew...
Updated 2 taps (homebrew/core and homebrew/cask).
==> New Formulae
[SNIP]
installer: The upgrade was successful.
🍺 xquartz was successfully installed!
quantum:~ mthomas$ which xquartz
/opt/X11/bin/xquartz
```

Step 2: Log on to TSCC, using **-Y** option (trusted)

```
[mthomas@home]$ ssh -Y tsccl-login.sdsc.edu -lmthomas
```

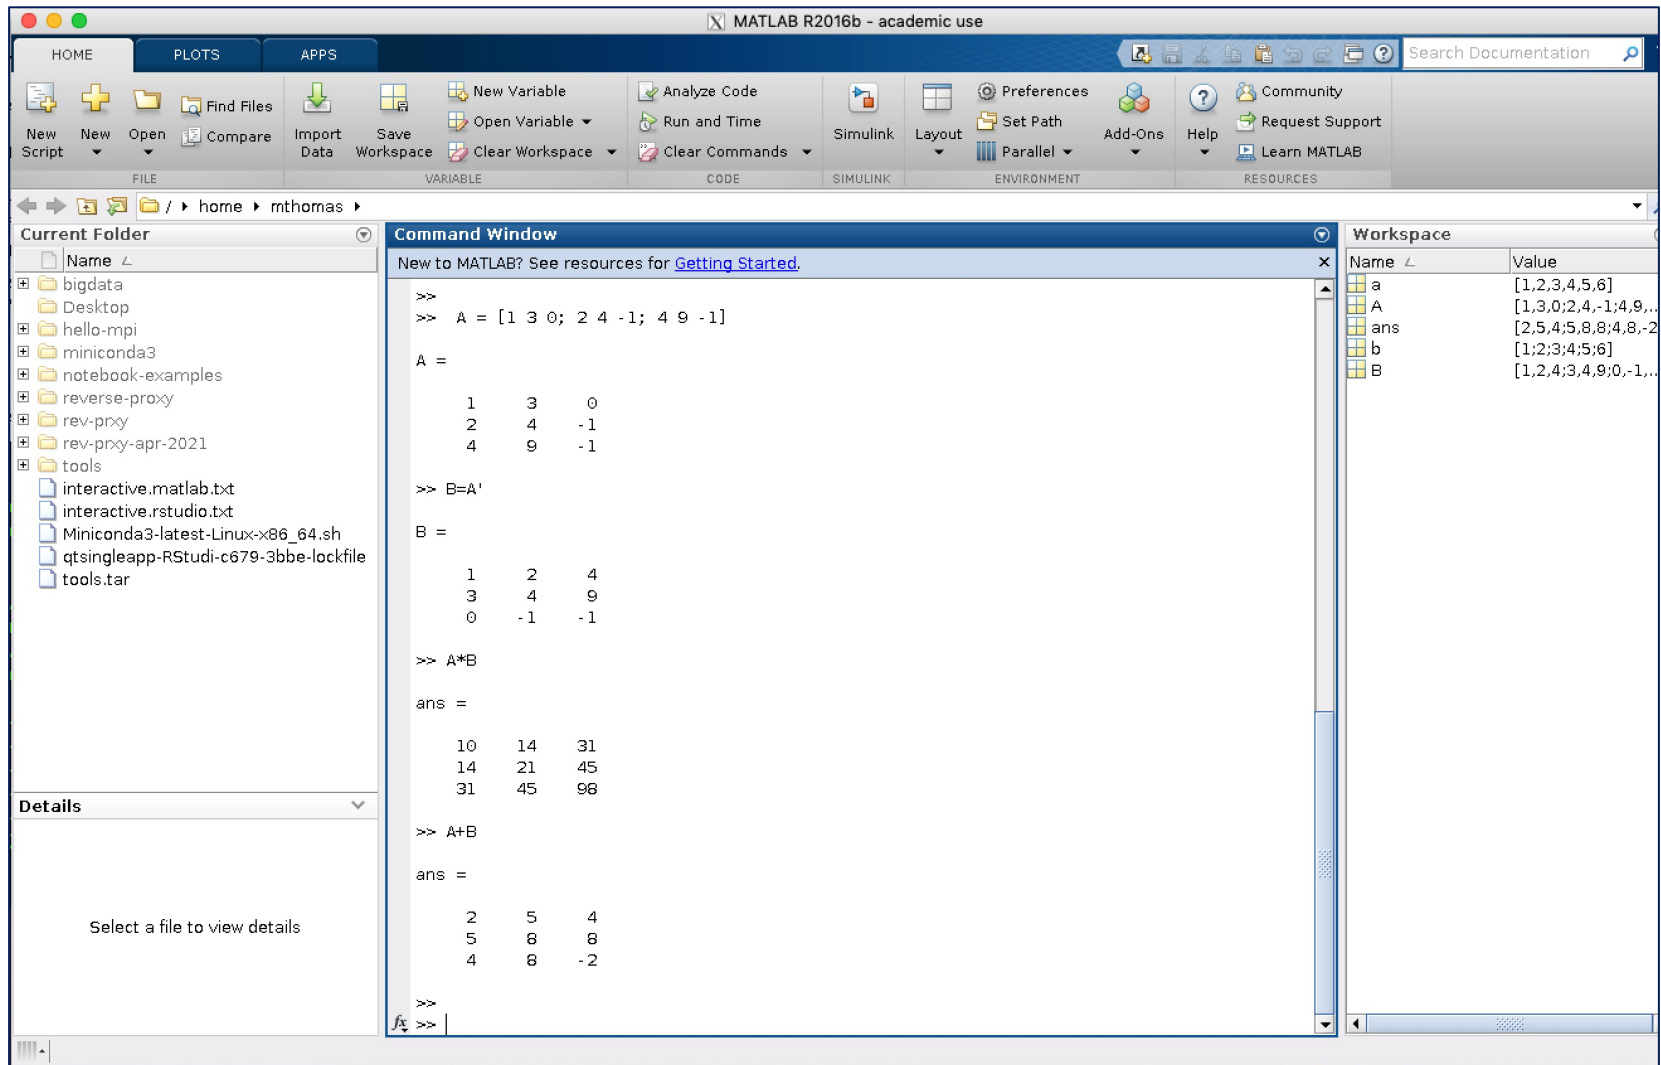
Step 3: Request an interactive node, using **-X** option

```
[mthomas@tsccl-login1 hello-mpi]$ qsub -I -X -l
nodes=1:ppn=1 -q glean
qsub: waiting for job 26843028.tsccl-mgr7.local to start
qsub: job 26843028.tsccl-mgr7.local ready
[mthomas@tsccl-4-46 ~]$
```

Step 4: Setup your module environment and run Matlab

```
[mthomas@tsccl-4-46 ~]$ module load matlab
[mthomas@tsccl-4-46 ~]$ matlab
MATLAB is selecting SOFTWARE OpenGL rendering.
```

TSCC Interactive Jobs: Running Matlab with GUI



TSCC Interactive Jobs: Running R console (no GUI)

Step 3: Request an interactive node, using **-X** option

```
[mthomas@tscc-login2 ~]$ qsub -l -q glean -l nodes=1:ppn=1
qsub: waiting for job 26844488.tscc-mgr7.local to start
qsub: job 26844488.tscc-mgr7.local ready
[mthomas@tscc-4-46 ~]$
```

Step 4: Setup your module environment and run Matlab

```
[mthomas@tscc-4-46 ~]$ module load R
Unloading compiler-dependent module openmpi_ib/3.1.4
[mthomas@tscc-4-46 ~]$ R
R version 4.0.2 (2020-06-22) -- "Taking Off Again"
Copyright (C) 2020 The R Foundation for Statistical Computing
Platform: x86_64-pc-linux-gnu (64-bit)
```

R is free software and comes with ABSOLUTELY NO WARRANTY.
[SNIP]

R is a collaborative project with many contributors.

[SNIP]

```
> myString <- "Hello, World!"
```

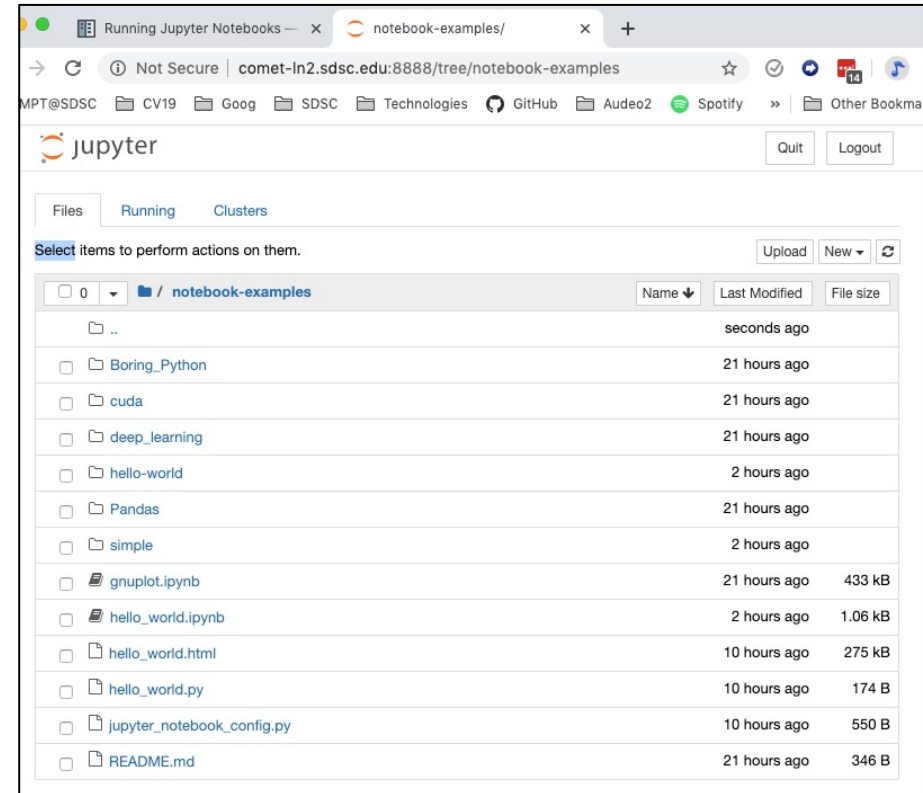
```
> print ( myString )
```

```
[1] "Hello, World!"
```

```
>
```

TSCC Interactive Jobs: Jupyter Notebooks

- What is Jupyter?
 - Free, open-source, interactive web tool known as a computational notebook
- Jupyter Services:
 - Notebooks (single user)
 - JupyterLab: advanced version of notebook
 - JupyterHub: multiuser.
- Jupyter Has a Key Vulnerability:
 - Jupyter Servers Provide Access to HPC File Systems
 - By default serve up as HTTP
- SDSC has created a solution:
 - **Satellite Reverse Proxy Service**
 - See presentation by Marty Kandes in this workshop



Resources

- TSCC User Guide
 - https://www.sdsc.edu/support/user_guides/tsc.html
- GitHub Repo for this workshop:
 - <https://github.com/sdsc-hpc-training-org/Using-Python-and-Jupyter-Notebooks-on-TSCC/>
- GitHub Repo for Example Notebooks:
 - <https://github.com/sdsc-hpc-training-org/notebook-examples>
- TSCC Support email: tsc-support@ucsd.edu
- TSCC Forum –
 - <https://mailman.ucsd.edu/mailman/listinfo/tsc-l>
 - Email: tsc-l@mailman.ucsd.edu
- SDSC Training Resources
 - https://www.sdsc.edu/education_and_training/training

Thank You