

galileo: A shell utility to help you launch Jupyter notebooks in a simple, secure way

Marty Kandes, Ph.D.

Computational & Data Science Research Specialist
High-Performance Computing User Services Group
San Diego Supercomputer Center
University of California, San Diego

TSCC Workshop
Thursday, September 2nd, 2021
2:30PM - 3:00PM PDT

Today's Session

- ▶ What is galyleo? How does it work?
- ▶ How do you get started with galyleo?
- ▶ What are the limitations of galyleo?
- ▶ Q&A
- ▶ Hands-on Session



What is galyleo? How does it work?

galyleo is a shell utility to help you launch Jupyter notebooks on high-performance computing (HPC) systems – like TSCC – in a simple, secure way.

It works with the Satellite reverse proxy service and a batch job scheduler like PBS/Torqueto provide each Jupyter notebook server you start with its own one-time, token-authenticated HTTPS connection between the compute resources of the HPC system the notebook server is running on and your web browser.

<https://github.com/mkandes/galyleo>

History & Design of galyleo

galyleo is a 2nd-generation shell utility developed to orchestrate a user's interaction with both Satellite and PBS/Torque to start a Jupyter session within a batch job.

Developed to replace the 1st-generation start-jupyter shell utility; to support the Expanse User Portal and the SDSC HPC User Services Group long-term

Key features in design:

- ▶ Recreate same interactions with Satellite service
- ▶ Increase flexibility for users to configure software environment; but also try to make it simpler for them to do themselves
- ▶ Batch job script is generated completely on-the-fly
- ▶ Command-line argument driven

Getting Started with galyleo in 3 Easy Steps

1. Prepend its installation location to your PATH.

```
export PATH="/projects/builder-group/galyleo:${PATH}"
```

2. launch your Jupyter notebook session.

```
galyleo launch --account abc123 --cpus 1 --time-limit 00:30:00  
--env-modules python,scipy/3.6 --jupyter notebook
```

3. Copy and paste the HTTPS URL into your web browser.

galyleo launch command-line options

```
galyleo launch --account abc123 --cpus 1 --time-limit 00:30:00 \  
--env-modules python,scipy/3.6 --jupyter notebook
```

A number of command-line options allow you to configure:

- ▶ the compute resources required to run your Jupyter notebook session;
- ▶ the type of Jupyter interface you want to use for the session and the location of the notebook working directory; and
- ▶ the software environment that contains the jupyter notebook server and the other software packages you want to work with during the session.

Please see documentation for complete list of command-line options:

<https://github.com/mkandes/galyleo>


Defining your software environment

--sif

--env-modules



--conda-env



Limitations

- ▶ galyleo is still a prototype under active development; expect things to change over time without warning
- ▶ galyleo currently only works if you use `/bin/bash` for your SHELL; support for other SHELLs will likely be supported in the future
- ▶ galyleo has only limited support for PBS/Torque; initial design was for Slurm; TSCC will likely transition to Slurm in the not-too-distant future

Please report issues on GitHub:

<https://github.com/mkandes/galyleo>

Questions?



TSCC Quick Reference Guide

```
export PATH="/projects/builder-group/galyleo:${PATH}"
```

```
galyleo launch --account abc123 --cpus 1 --time-limit 00:30:00 \  
--env-modules python,scipy/3.6 --jupyter notebook
```

```
galyleo launch --account abc123 --cpus 1 --time-limit 00:30:00 \  
--sif r-notebook_latest.sif
```

```
export SIF_DIR='/projects/builder-group/singularity'  
galyleo launch --account abc123 --partition gpu-hotel --cpus 3 \  
--time-limit 00:30:00 \  
--sif "${SIF_DIR}/pytorch/pytorch-v1.4.0-gpu-20200224.simg" \  
--bind '/oasis,${TMPDIR}"' --nv
```

```
galyleo launch --account abc123 --cpus 1 --time-limit 00:30:00 \  
--conda-env notebooks-sharing \  
--conda-init miniconda3/etc/profile.d/conda.sh
```

```
galyleo launch --account abc123 --cpus 1 --time-limit 00:30:00 \  
--conda-env notebooks-sharing \  
--conda-pack notebooks-sharing.tar.gz \  
--scratch-dir "${TMPDIR}"
```